



NATIONAL RESEARCH
UNIVERSITY

Faculty of Computer Science
School of Software Engineering

SQLITE RDBMS EXTENSION FOR DATA INDEXING USING B-TREE MODIFICATIONS

Student: Anton Rigin

Academic Supervisor: Sergey Shershakov,
Senior Lecturer at School of Software Engineering

Presentation at SYRCoSE 2019

Saratov, 2019


OUTLINE

- B-tree and modifications
- SQLite and extensions
- B-tree modifications C++ library and research conducted
- Motivation, existing solutions and main goals
- Algorithm of the best index structure
- Technologies used
- Implementation
- Usage example and experiment conducted

B-TREE

- Balanced search tree [1]
- Nodes may contain more than 1 key and more than 2 pointers to the children nodes [1]
- If some node contains k keys than it contains $k + 1$ pointers to the children nodes [1]
- **B-tree order** is such a t number that (k is the count of keys in the node):
 - ✓ for each non-root node: $t - 1 \leq k \leq 2t - 1$ [1]
 - ✓ for root node in the non-empty tree: $1 \leq k \leq 2t - 1$ [1]
 - ✓ for root node in the empty tree: $k = 0$ [1]
- **B-tree height** is $O(\log_t n)$, where t is tree order and n is the count of keys in the tree [1]
- Usually used as the data index [1]

B-TREE



The B-tree example, $t = 6$

B-TREE OPERATIONS

- **Searching**

- ✓ Time complexity – $O(t \log_t n)$ [1]
- ✓ Memory usage – $O(t)$ [1]
- ✓ Disk operations count – $O(\log_t n)$ [1]

- **Nodes split (the part of insertion)**

- ✓ Time complexity – $O(t)$ [1]
- ✓ Memory usage – $O(t)$ [1]
- ✓ Disk operations count – $O(1)$ [1]

- **Insertion**

- ✓ Time complexity – $O(t \log_t n)$ [1]
- ✓ Memory usage – $O(t \log_t n)$ for simple recursion and $O(t)$ for tail recursion or loop [1]
- ✓ Disk operations count – $O(\log_t n)$ [1]

- **Deletion**

- ✓ Time complexity – $O(t \log_t n)$ [1]
- ✓ Memory usage – $O(t \log_t n)$ for simple recursion and $O(t)$ for tail recursion or loop [1]
- ✓ Disk operations count – $O(\log_t n)$ [1]

B-TREE MODIFICATIONS

- **B⁺-tree**
 - ✓ Only leaf nodes contain real keys (real data), other nodes contain router keys [2]
 - ✓ Deletion is probably faster than in B-tree
- **B^{*}-tree**
 - ✓ Each node (except of the root node) is filled at least by 2/3 not 1/2 [3]
 - ✓ Keys insertion in B^{*}-tree is expected to be faster than in B-tree
- **B⁺⁺-tree**
 - ✓ Developed previously by author of this work [4]
 - ✓ Combines the main B⁺-tree and B^{*}-tree features together [4]

[2] K. Pollari-Malmi. (2010). B+-trees [PDF paper]. Available: <https://www.cs.helsinki.fi/u/mluukkai/tirak2010/B-tree.pdf>

[3] "B^{*}-tree." NIST Dictionary of Algorithms and Data Structures. Available: <https://xlinux.nist.gov/dads/HTML/bstartree.html> (accessed Dec. 24, 2018).

[4] A. Rigin, "On the Performance of Multiway Trees in the Problem of Structured Data Indexing," (in Russian), coursework, Dept. Soft. Eng., HSE, Moscow, Russia, 2018. 6


SQLITE

- Popular open-source embedded relational DBMS
- Written in the C language
- Uses the B-tree as the default index
- SQLite extensions are the dynamically linked libraries [5]





B-TREE MODIFICATIONS C++ LIBRARY


- Developed previously [4]
- Contains B-tree, B⁺-tree, B^{*}-tree and B⁺⁺-tree implementations
- In the current work connected to the SQLite as the **run-time loadable extension**





RESEARCH CONDUCTED USING THE B-TREE MODIFICATIONS C++ LIBRARY




RESEARCH CONDUCTED USING THE B-TREE MODIFICATIONS C++ LIBRARY



RESEARCH CONDUCTED USING THE B-TREE MODIFICATIONS C++ LIBRARY



RESEARCH CONDUCTED USING THE B-TREE MODIFICATIONS C++ LIBRARY



MOTIVATION AND EXISTING SOLUTIONS

- **Motivation**

- ✓ In the current time, the processed data amounts are continuously growing
- ✓ It is necessary to develop new efficient approaches to data indexing in the DBMSs
- ✓ SQLite contains a small number of index structures by default – it is relevant to add new ones

- **Existing solutions**

- ✓ B-tree is the default index structure in the SQLite
- ✓ There are some SQLite extensions (for example, the extension adding the R-tree indexing)
- ✓ No extensions with B⁺-tree, B^{*}-tree, and B⁺^{*}-tree were found

THE MAIN GOALS OF THE WORK

- To add B-tree modifications such as B^+ -tree, B^* -tree and B^{*+} -tree to SQLite
- To develop and implement an algorithm that would allow selecting the most appropriate indexing data structure (B -tree, B^+ -tree, B^* -tree or B^{*+} -tree) when a user manipulates a table


ALGORITHM OF SELECTING THE BEST INDEX STRUCTURE

- Selects from the B-tree and its modifications (B^+ -tree, B^* -tree and B^{*+} -tree)
- Runs at the start of each table operation (search, insertion, updating, deletion)
- Performs the index structure rebuilding only on each 1000-th operation and only for the first 10000 operations
- Index structure selection depends on the ratios of the numbers of operations of different types (search, insert, delete) on the tree
- The best index structure is defined by the step 6 of the algorithm
- The tree order of the B-trees and their modifications used in the SQLite extension developed in this work equals 100

ALGORITHM OF SELECTING THE BEST INDEX STRUCTURE

1. If the current total count of the operations on the tree is equal to 0 or more than 10000 or not a multiple of 1000, then exit the algorithm, otherwise go to step 2.
2. If the current count of the modifying operations on the tree (key insertions, key deletions) is less than 10 % of the current total count of the operations on the tree, then exit the algorithm, otherwise go to step 3.
3. If the current count of the key insertion operations is more than 90 % of the current count of the modifying operations on the tree, then select the B*-tree as the index structure and go to step 6, otherwise go to step 4.
4. If the current count of the key insertion operations is not less than 60 % of the current count of the modifying operations on the tree, then select the B⁺-tree as the index structure and go to step 6, otherwise go to step 5.
5. Select the B⁺-tree as the index structure and go to step 6.
6. If the new index structure was selected in the steps 3 – 5, then rebuild the existing index structure into the new selected index structure (which is considered as the best index structure) saving all the data stored in the existing index structure.

TREE KEY STRUCTURE



TECHNOLOGIES USED


THE
C
PROGRAMMING
LANGUAGE



C++



SQLite
Studio



IMPLEMENTATION

- The C API for the B-tree modifications library is implemented as dynamically shared library using the *extern "C" { ... }* C++ statement
- The SQLite extension registers the *btrees_mods* module for creating the virtual tables
- Virtual table is any table created using such a module
- The extension also registers the *btreesModsVisualize*, *btreesModsGetTreeOrder*, *btreesModsGetTreeType* functions

IMPLEMENTATION

Method	Purpose
btreesModsCreate (sqlite3*, void*, int, const char* const*, sqlite3_vtab**, char**)	Creates a new table.
btreesModsUpdate (sqlite3_vtab*, int, sqlite3_value**, sqlite_int64*)	Inserts, deletes or updates a value of a row in the table.
btreesModsFilter (sqlite3_vtab_cursor*, int, const char*, int, sqlite3_value**)	Searches for a row in the table.

IMPLEMENTATION

Method	Purpose
btreesModsVisualize (sqlite3_context*, int, sqlite3_value**)	Outputs the graphical representation of the table's index structure (tree) into the GraphViz DOT file.
btreesModsGetTreeOrder (sqlite3_context*, int, sqlite3_value**)	Outputs the order of the tree used as the table's index structure.
btreesModsGetTreeType (sqlite3_context*, int, sqlite3_value**)	Outputs the type of the tree (1 – B-tree, 2 – B ⁺ -tree, 3 – B [*] -tree, 4 – B ⁺ [*] -tree) used as the table's index structure.

USAGE EXAMPLE

```
SQLite version 3.26.0 2018-12-01 12:34:55
Enter ".help" for usage hints.
Connected to a transient in-memory database.
Use ".open FILENAME" to reopen on a persistent database.
sqlite> .load ./btrees_mods
sqlite> CREATE VIRTUAL TABLE btt USING btrees_mods(id INTEGER PRIMARY KEY, a INTEGER, b TEXT);
sqlite> INSERT INTO btt VALUES (4, 2, "ABC123");
sqlite> INSERT INTO btt VALUES (7, 3, "def");
sqlite> SELECT * FROM btt WHERE id = 4;
4|2|ABC123
sqlite> SELECT * FROM btt WHERE id = 7;
7|3|def
sqlite> SELECT * FROM btt WHERE id = 4 OR id = 7;
4|2|ABC123
7|3|def
sqlite> .tables
btrees_mods_idxinfo  btt                                btt_real
sqlite> SELECT * FROM btt_real;
4|2|ABC123
7|3|def
sqlite> SELECT * FROM btrees_mods_idxinfo;
btt|1|0|id|INTEGER|4|tree_18291557263097.btree
sqlite> DROP TABLE btt;
sqlite> .tables
btrees_mods_idxinfo
sqlite> SELECT * FROM btrees_mods_idxinfo;
sqlite> .exit
```

EXPERIMENT CONDUCTED USING THE SQLITE EXTENSION

Operation on the table	Total execution time (ms)	Mean execution time per row (ms)
Table creation	21	-
First 500 rows insertion	9128	18.3
Next 500 rows insertion	9802	19.6
1001st row insertion (including the B-tree into the B*-tree rebuilding)	50	50
Next 499 rows insertion	9168	18.4
Last 500 rows insertion	8933	17.9

EXPERIMENT CONDUCTED USING THE SQLITE EXTENSION

Operation on the table	Total execution time (ms)	Mean execution time per row (ms)
First 500 rows deletion	9888	19.8
Next 500 rows deletion (including the B*-tree into the B ⁺ -tree rebuilding)	9974	19.9
Next 500 rows deletion	9784	19.6
Last 500 rows deletion (including the B ⁺ -tree into the B ⁺ -tree rebuilding)	9201	18.4
1000 rows insertion	18113	18.1
Next 4800 rows insertion (including the B ⁺ -tree into the B ⁺ -tree rebuilding)	86465	18

EXPERIMENT CONDUCTED USING THE SQLITE EXTENSION

- The key insertion on the B^* -tree is faster than on the B-tree as expected
- The key deletion on the B^+ -tree and B^{*+} -tree is faster than on the B-tree and B^* -tree respectively
- The key insertion on the B^{*+} -tree is slightly faster than on the B^+ -tree

SUMMARY

- B-trees modifications library is connected to the SQLite as an extension using developed in this work C API
- Research conducted using this library is presented
- Algorithm of selecting the best index structure is developed and implemented
- Experiment is conducted using the SQLite extension

REFERENCES

- [1] R. Bayer and E. McCreight, “Organization and maintenance of large ordered indexes,” *Acta Informatica*, vol. 1, no. 3, pp. 173 – 189, 1972.
- [2] K. Pollari-Malmi. (2010). “B⁺-trees,” *University of Helsinki*. [PDF paper]. Available: <https://www.cs.helsinki.fi/u/mluukkai/tirak2010/B-tree.pdf>
- [3] “B*-tree.” NIST Dictionary of Algorithms and Data Structures. Available: <https://xlinux.nist.gov/dads/HTML/bstartree.html> (accessed Dec. 24, 2018).
- [4] A. Rigin, “On the Performance of Multiway Trees in the Problem of Structured Data Indexing,” (in Russian), coursework, Dept. Soft. Eng., HSE, Moscow, Russia, 2018.
- [5] “Run-Time Loadable Extensions.” SQLite.org. Available: <https://www.sqlite.org/loadext.html> (accessed Jan. 20, 2019).



NATIONAL RESEARCH
UNIVERSITY

Thank you for your attention!

amrigin@edu.hse.ru

anton19979@yandex.ru

anton19979@yandex-team.ru